

Practical 1:-Introduction to R

```
# Lab 1

# Print a message
print('Hello, Welcome to R Programming')

#Assign variables
x <- 25
y <- 5

#Arithmetic operations
sum_val <- x + y
diff_val <- x - y
prod_val <- x * y
div_val <- x / y

sum_val
diff_val
prod_val
div_val

# Relational operations
greater_check <- x > y
equal_check <- x == y

# Load and explore iris dataset
data("iris") #load dataset
head(iris) #first 6 rows
str(iris) #structure of dataset
summary(iris) #descriptive statistics
```

Practical 2:-Importing and Exploring a Dataset in R

```
# Lab 2

#Load required package
library(dplyr)

#1. Import inbuilt dataset: iris
data("iris")

head(iris) #first 6 rows

str(iris) # structure

summary(iris) #summary statistics

#2. Import another dataset: mtcars

data("mtcars")

head (mtcars)

str(mtcars)

#3. Titanic dataset

data("Titanic")

Titanic #view dataset

#4. Check for missing values

sum(is.na(iris)) # total missing values

colSums(is.na(mtcars)) #missing values per column

#5. Summarize distribution by group

iris %>%
  group_by(Species) %>%
  summarise(
    Avg_Sepal_Length = mean(Sepal.Length),
    Avg_Petal_Length = mean(Petal.Length)
  )
```

Practical 3:-Data cleaning and Preprocessing in R

```
# Lab 3

library(dplyr)

# Load dataset

data("airquality")

# Display the first few rows of the original dataset

print("Original Airquality Dataset Head:")

head(airquality)

# 1. Check for missing values

print("--- Missing Value Check (airquality) ---")

# Total NA count across the entire dataset

print(paste("Total NA count:", sum(is.na(airquality)))))

# Missing values per column

print("Missing values per column:")

colSums(is.na(airquality))

# 2. Handle missing values

print("--- Handling Missing Values (airquality) ---")

# Replace missing Ozone values with the mean of the column

# 'na.rm = TRUE' ensures the mean is calculated by ignoring existing NA values

airquality$Ozone[is.na(airquality$Ozone)] <- mean(airquality$Ozone, na.rm = TRUE)

# Replace missing Solar.R values with the median of the column

airquality$Solar.R[is.na(airquality$Solar.R)] <- median(airquality$Solar.R, na.rm = TRUE)

# Verify the changes

print("Missing values per column after imputation:")

colSums(is.na(airquality))

# 3. Remove duplicates (using iris dataset as an example)

data("iris")

print("--- Duplicates Removal (iris) ---")
```

```
# Creates a clean dataset by selecting rows that are NOT duplicates.  
# 'duplicated()' returns TRUE for the second and subsequent occurrences of duplicate rows.  
iris_clean <- iris[!duplicated(iris), ]  
  
# Display the number of rows before and after cleaning  
print(paste("Original rows in iris:", nrow(iris)))  
print(paste("Rows after removing duplicates:", nrow(iris_clean)))  
  
# 4. Standardize and use formats  
  
print("--- Data Standardization (iris) ---")  
  
# Convert the Species factor to character, then to lowercase for standardization  
iris_clean$Species <- tolower(as.character(iris_clean$Species))  
  
# Convert the standardized character column back to a factor  
iris_clean$Species <- as.factor(iris_clean$Species)  
  
# View the cleaned dataset head  
print("Cleaned Iris Dataset Head:")  
head(iris_clean)
```

Practical 4:-Description Statistics and Basic Visualization in R

```
# Lab 4

library(dplyr)
library(ggplot2)
library(modeest) # for mode

# Load dataset
data("iris")

# Descriptive statistics
mean(iris$Sepal.Length) # mean
median(iris$Sepal.Length) # median
mlv(iris$Sepal.Length, method="mfv") # mode
sd(iris$Sepal.Length) # standard deviation
range(iris$Sepal.Length) # min and max
summary(iris$Sepal.Length) # summary

# Histogram
hist(iris$Sepal.Length,
     main="Histogram of Sepal Length",
     xlab="Sepal Length", col="lightblue", border="black")

# Scatterplot
plot(iris$Sepal.Length, iris$Petal.Length,
      main="Scatterplot of Sepal vs Petal Length",
      xlab="Sepal Length", ylab="Petal Length",
      col=c("blue", pch=19))

# Boxplot
boxplot(Sepal.Length ~ Species, data=iris,
        main="Boxplot of Sepal Length by Species",
        xlab="Species", ylab="Sepal Length",
        col=c("lightgreen", "lightblue", "pink"))
```

Prcatical 5:-Variable Transformation and Feature Engineering in R

```
# Lab 5

library(dplyr)
library(caret)

# Load dataset
data("mtcars")

# 1. Binning: Categorise mpg into Low, Medium, High
mtcars$mpg_category = cut(mtcars$mpg,
                           breaks = c(-Inf, 15, 25, Inf),
                           labels = c("Low", "Medium", "High"))

table(mtcars$mpg_category)

# 2. Encoding: Convert Species to numeric codes (iris dataset)
data("iris")

iris$Species_code = as.numeric(as.factor(iris$Species))

head(iris[, c("Species", "Species_code")])

# 3. Normalization: Scale wt (weight) column
mtcars$wt_normalized = (mtcars$wt - min(mtcars$wt)) /
  (max(mtcars$wt) - min(mtcars$wt))

head(mtcars$wt_normalized)

# 4. Standardization: Z-score for hp (horsepower)
mtcars$hp_zscore = scale(mtcars$hp)

head(mtcars$hp_zscore)

# 5. Feature Creation: Power-to-Weight Ratio
mtcars$power_to_weight = mtcars$hp / mtcars$wt

head(mtcars$power_to_weight)
```

Practical 6:-Exploratory Data Analysis (EDA) in R

```
# Lab 6

library(dplyr)
library(ggplot2)
library(GGally)
library(corrplot)
library(ggcormplot)

# Load dataset

data("iris")

head(iris)

# 1. Summary statistics

summary(iris)

# 2. Histogram of Sepal.Length

ggplot(iris, aes(x = Sepal.Length)) +
  geom_histogram(bins = 15, fill = "lightblue", color = "black") +
  labs(title = "Distribution of Sepal Length", x = "Sepal Length", y = "Frequency")

# 3. Scatterplot Sepal Length vs Petal Length

ggplot(iris, aes(x = Sepal.Length, y = Petal.Length, color = Species)) +
  geom_point(size = 3) +
  labs(title = "Sepal Length vs Petal Length")

# 4. Boxplot of Sepal Width by Species

ggplot(iris, aes(x = Species, y = Sepal.Width, fill = Species)) +
  geom_boxplot() +
  labs(title = "Boxplot of Sepal Width by Species")

# 5. Correlation matrix

corr_matrix = cor(iris[, 1:4])

corr_matrix

ggcormplot(corr_matrix, lab = TRUE, title="Correlation Matrix Heatmap")

# 6. Pair plot

ggpairs(iris[, 1:4])
```

Practical 7:-Statistical tests In R

```
# Lab 7

library(dplyr)

# Load iris dataset

data("iris")

# 1. t-test: Compare Sepal.Length of setosa and versicolor

t_test_result <- t.test(Sepal.Length ~ Species,
                        data = iris %>% filter(Species %in% c("setosa",
                        "versicolor")))

t_test_result

# 2. ANOVA: Compare Sepal.Length across all species

anova_model <- aov(Sepal.Length ~ Species, data = iris)

summary(anova_model)

# 3. Correlation: Sepal Length and Petal Length

correlation <- cor(iris$Sepal.Length, iris$Petal.Length)

correlation

# 4. Correlation test with significance

cor_test <- cor.test(iris$Sepal.Length, iris$Petal.Length)

cor_test
```

Practical 8:-Regression Analysis in R

```
# Lab 8

library(dplyr)
library(ggplot2)

# Load dataset
data("mtcars")
head(mtcars)

# 1. Simple Linear Regression: mpg predicted by wt
model_simple = lm(mpg ~ wt, data = mtcars)
summary(model_simple)

# 2. Plot regression line
ggplot(mtcars, aes(x = wt, y = mpg)) +
  geom_point(color = "blue") +
  geom_smooth(method = "lm", se = TRUE, color = "red") +
  labs(title = "Simple Linear Regression: MPG vs Weight",
       x = "Weight (1000 lbs)", y = "Miles Per Gallon")

# 3. Multiple Linear Regression: mpg predicted by wt and hp
model_multiple = lm(mpg ~ wt + hp, data = mtcars)
summary(model_multiple)

# 4. Residual diagnostics
plot(model_multiple, which = 1) # Residuals vs Fitted
plot(model_multiple, which = 2) # Q-Q Plot
```

Practical 9:-Classification Analysis in R

```
library(dplyr)
```

```
library(caret)
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
library(ROCR)
```

```
library(ggplot2)
```

```
data("iris")
```

```
iris_bin <- iris %>%
```

```
  mutate(Species = factor(ifelse(Species == "setosa", "setosa", "non_setosa"),  
    levels = c("non_setosa", "setosa")))
```

```
set.seed(123)
```

```
idx <- createDataPartition(iris_bin$Species, p = 0.7, list = FALSE)
```

```
trainData <- iris_bin[idx, ]
```

```
testData <- iris_bin[-idx, ]
```

```
log_model <- glm(Species ~ Sepal.Length + Petal.Length,
```

```
  data = trainData, family = binomial())
```

```
log_prob <- predict(log_model, testData, type = "response")
```

```
log_class <- ifelse(log_prob > 0.5, "setosa", "non_setosa")
```

```
confusionMatrix(factor(log_class, levels = levels(testData$Species)),
```

```
  testData$Species)
```

```
actual_labels <- ifelse(testData$Species == "setosa", 1, 0)
```

```
pred_obj <- prediction(log_prob, actual_labels)
```

```
perf <- performance(pred_obj, "tpr", "fpr")
plot(perf, col = "blue", main = "ROC Curve - Logistic Regression")
abline(a = 0, b = 1, lty = 2, col = "red")

tree_model <- rpart(Species ~ Sepal.Length + Petal.Length,
                     data = trainData, method = "class")
rpart.plot(tree_model)
tree_pred <- predict(tree_model, testData, type = "class")
confusionMatrix(tree_pred, testData$Species)
```

Practical 10:-Clustering Analysis in R(K-Means)

```
# Lab 10

library(dplyr)
library(ggplot2)
library(cluster)
library(factoextra)

# Load dataset (only numeric features)
data("iris")

iris_data <- iris[, 1:4]

# 1. Elbow Method to find optimal K
fviz_nbclust(iris_data, kmeans, method = "wss") +
  labs(title = "Elbow Method for Optimal K")

# 2. Apply K-means clustering with K=3
set.seed(123)

kmeans_model <- kmeans(iris_data, centers = 3, nstart = 20)

# 3. Cluster assignments
kmeans_model$cluster[1:10]
table(kmeans_model$cluster, iris$Species)

# 4. Visualize clusters
fviz_cluster(kmeans_model, data = iris_data,
  ellipse.type = "norm",
  palette = "jco",
  ggtheme = theme_minimal())
```